

Verification of Data Layout Transformations

Ramon Fernández Mir

with Arthur Charguéraud

Inria

17/09/2018

Motivating example

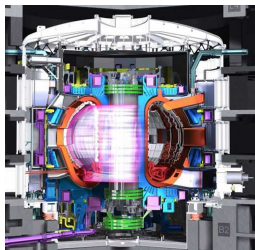


Figure: ITER tokamak

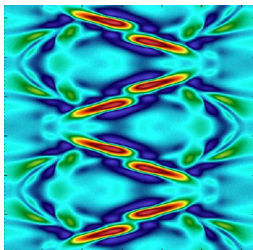


Figure: Plasma physics

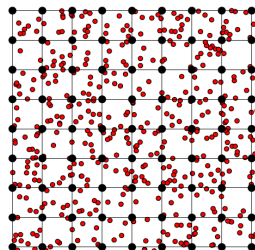


Figure: PIC simulation

Challenges:

- Exploit data-level parallelism.
- Use domain-specific knowledge of the code.
- Do it without introducing any bugs.

Motivating example - initial code

```
typedef struct {  
    // Position  
    float x, y, z;  
    // Other fields  
    float vx, vy, vz, c, m, v;  
} particle;  
  
particle data[NUM_PARTICLES];  
  
for (int i = 0; i < NUM_PARTICLES; i++) {  
    // Some calculation  
}
```

Motivating example - splitting

Suppose that the calculation uses mainly the position.

```
typedef struct {  
    float vx, vy, vz, c, m, v;  
} cold_fields;  
  
typedef struct {  
    float x, y, z;  
    cold_fields *other;  
} particle;  
  
particle data[NUM_PARTICLES];
```

Motivating example - peeling

Further suppose that the initial 'particle' record is not used as part of a dynamic data structure.

```
typedef struct {  
    float vx, vy, vz, c, m, v;  
} cold_fields;
```

```
typedef struct {  
    float x, y, z;  
} hot_fields;
```

```
cold_fields other_data[NUM_PARTICLES];  
hot_fields pos_data[NUM_PARTICLES];
```

Motivating example - AoS to SoA

Now, say that we want to take advantage of vector instructions.

```
typedef struct {  
    float x[NUM_PARTICLES];  
    float y[NUM_PARTICLES];  
    float z[NUM_PARTICLES];  
} hot_fields;  
  
hot_fields pos_data;
```

Motivating example - AoS to AoSoA

But without reducing too much the locality between accesses to fields of the original struct.

```
typedef struct {  
    float x[N];  
    float y[N];  
    float z[N];  
} hot_fields;
```

```
hot_fields pos_data[NUM_PARTICLES / N];
```

Motivating example - summary

In short, the transformations we have seen are:

- Splitting.
- Peeling.
- AoS to SoA.
- AoS to AoSoA.

Note that after all these changes, where we wrote:

```
data[i].x
```

Now we have to write:

```
pos_data[i / N].x[i % N]
```


Project goals

- Find the basic transformations that combined give rise to the ones we are interested in.
- Formalize a C-like language with arrays, structs and pointers.
 - On a high-level, to simplify the proofs.
 - On a low-level, to be closer to the semantics of C.
- Define the transformations and prove their correctness.

Basic transformations

1. Field grouping

```
// Before
typedef struct {
    int a, b, c;
} s;
```

```
// After
typedef struct {
    int b, c;
} sg;
```

```
typedef struct {
    int a; sg fg;
} s';
```

2. Array tiling

```
// Before
int a[N];
```

```
// After
int a'[N/_B][_B][B];
```

3. Adding indirection

```
// Before
typedef struct {
    int a, b;
} s;
```

```
// After
typedef struct {
    int a; int *b;
} s';
```

4. AoS to SoA


```
// Before
typedef struct {
    int a, b;
} s;
```

```
// After
typedef struct {
    int a[N]; int b[N];
} s;
```

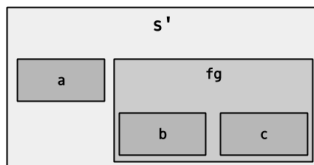
Basic transformations - grouping

1. Field grouping

```
// Before  
typedef struct {  
    int a, b, c;  
} s;
```



```
// After  
typedef struct {  
    int b, c;  
} sg;  
  
typedef struct {  
    int a; sg fg;  
} s';
```



Basic transformations - tiling

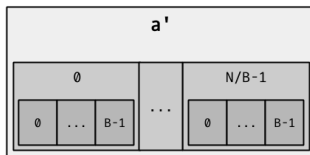
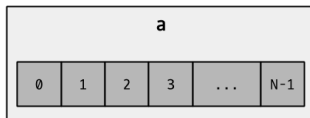
2. Array tiling

```
// Before
```

```
int a[N];
```

```
// After
```

```
int a'[N/_B][_B];
```

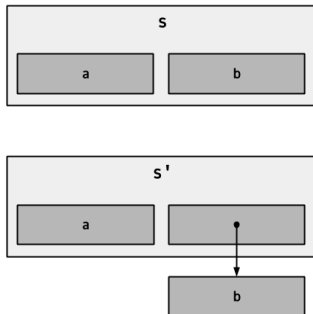


Basic transformations - indirection

3. Adding indirection

```
// Before  
typedef struct {  
    int a, b;  
} s;
```

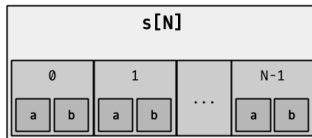
```
// After  
typedef struct {  
    int a; int *b;  
} s';
```



Basic transformations - AoS to SoA

4. AoS to SoA

```
// Before  
typedef struct {  
    int a, b;  
} s;
```



```
// After  
typedef struct {  
    int a[N]; int b[N];  
} s';
```



Basic transformations - justification

- **Peeling:** Field grouping twice.
- **Splitting:** Field grouping and then adding indirection on the field holding the group.
- **AoS to SoA:** AoS to SoA.
- **AoS to AoSoA:** Array tiling and then AoS to SoA on the tiles.

Language overview - values and terms

Inductive val : Type :=

- | val_error : val
- | val_unit : val
- | val_uninitialized : val
- | val_bool : bool → val
- | val_int : int → val
- | val_double : int → val
- | val_abstract_ptr : loc → accesses → val
- | val_array : typ → list val → val
- | val_struct : typ → map field val → val

Inductive trm : Type :=

- | trm_var : var → trm
- | trm_val : val → trm
- | trm_if : trm → trm → trm → trm
- | trm_let : bind → trm → trm → trm
- | trm_app : prim → list trm → trm
- | trm_while : trm → trm → trm
- | trm_for : var → val → val → trm → trm.

Language overview - primitive operations

```
Inductive prim : Type :=  
| prim_binop : binop → prim  
| prim_get : typ → prim  
| prim_set : typ → prim  
| prim_new : typ → prim  
| prim_new_array : typ → prim  
| prim_struct_access : typ → field → prim  
| prim_array_access : typ → prim  
| prim_struct_get : typ → field → prim  
| prim_array_get : typ → prim
```

Examples of the semantics of our language compared to C:

get p : *p	array_access p i : p + i
set p v : *p = v	struct_access p f : &(p->f)
new T : malloc(sizeof(T))	struct_get s f : s.f

where pointers are represented as pairs:

(l, [access_field T f, access_array T' i])

which would correspond to the address:

$l + \text{field_offset}(f) + i * \text{sizeof}(T')$

Language overview - semantics

Some crucial definitions:

Definition $\text{typdefctx} := \text{map typvar typ}.$

$\text{Record ll_typdefctx} := \text{make_ll_typdefctx} \{$
 $\text{typvar_sizes} \quad : \text{map typvar size};$
 $\text{fields_offsets} : \text{map typvar (map field offset)};$
 $\text{fields_order} \quad : \text{map typvar (list field)} \}.$

Definition $\text{stack} := \text{Ctx.ctx val}.$

Definition $\text{state} := \text{map loc val}.$

And the relation that defines the big-step reduction rules:

$$\text{red} \subseteq \text{typdefctx} \times \text{ll_typdefctx} \times \text{stack} \times \text{state} \times \text{trm} \times \text{state} \times \text{val}$$

Language overview - typing

The allowed types are:

```
Inductive typ : Type :=  
  | typ_unit : typ  
  | typ_int : typ  
  | typ_double : typ  
  | typ_bool : typ  
  | typ_ptr : typ → typ  
  | typ_array : typ → option size → typ  
  | typ_struct : map field typ → typ  
  | typ_var : typvar → typ.
```

With their corresponding definitions (analogous to stack and state):

```
Definition gamma : Ctx.ctx typ.
```

```
Definition phi : map loc typ.
```

Typing is defined as the following relation:

$$\text{typing} \subseteq \text{typdefctx} \times \text{gamma} \times \text{phi} \times \text{trm} \times \text{typ}$$

Language overview - properties

An approximation to type safety:

Theorem `type_soundness` : $\forall C \text{ LLC } m \ t \ v \ T,$
 `red C LLC nil empty t m v` \rightarrow
 `typing C nil empty t T` \rightarrow
 $\exists f, \text{typing_val } C \ f \ v \ T$
 $\wedge \text{state_typing } C \ f \ m.$

TODO: Need to think of something...

Field grouping

The arguments of the transformation are:

- The name of the struct being changed.
- The fields being grouped.
- The name of the new struct that will contain said fields.
- The new field holding the new struct.

These are used to define a transformation for:

- type definitions contexts,
- terms,
- accesses,
- states and
- values,
- stacks.

Field grouping - OK

We first need a way of checking that the transformation is well-defined.

```
Inductive group_tr_ok : group_tr → typdefctx → Prop :=
| group_tr_ok_intros : ∀Tfs Tt fs fg Tg gt C,
  gt = make_group_tr Tt fs Tg fg →
  Tt ∈ dom C →
  (* The struct Tt can be transformed. *)
  C[Tt] = typ_struct Tfs →
  Tg ∉ dom C →
  fs ⊆ dom Tfs →
  fg ∉ dom Tfs →
  (* Tt doesn't appear anywhere else in the typdefctx. *)
  (∀ Tv,
    Tv ∈ dom C →
    Tv ≠ Tt →
    ~free_typvar C Tt C[Tv]) →
  group_tr_ok gt C.
```

Field grouping - typdefctx

We 'update' the type definitions context as follows:

```
Inductive tr_typdefctx (gt:group_tr) : typdefctx → typdefctx → Prop :=  
  | tr_typdefctx_intro : ∀Tfs Tfs' Tfsg Tt fs Tg fg C C',  
    gt = make_group_tr Tt fs Tg fg →  
    dom C' = dom C ∪ {Tg} →  
    (* The original map from fields to types. *)  
    C[Tt] = typ_struct Tfs →  
    (* The map for the new struct and for the grouped fields. *)  
    tr_struct_map gt Tfs Tfs' Tfsg →  
    C'[Tt] = typ_struct Tfs' →  
    C'[Tg] = typ_struct Tfsg →  
    (* The other type variables stay the same. *)  
    (∀ T,  
      T ∈ dom C →  
      T ≠ Tt →  
      C'[T] = C[T]) →  
    tr_typdefctx gt C C'.
```

Field grouping - accesses

For accesses, if we look at the interesting case:

```
Inductive tr_accesses (gt:group_tr) : accesses → accesses → Prop :=  
  | tr_accesses_field_group : ∀Tt fs fg Tg f a0 p a1 a2 p',  
    gt = make_group_tr Tt fs Tg fg →  
    f ∈ fs →  
    (* The access s.f *)  
    a0 = access_field (typ_var Tt) f →  
    (* Becomes s'.fg.f *)  
    a1 = access_field (typ_var Tt) fg →  
    a2 = access_field (typ_var Tg) f →  
    tr_accesses gt p p' →  
    tr_accesses gt (a0::p) (a1::a2::p')
```

This is used in:

```
Inductive tr_val (gt:group_tr) : val → val → Prop :=  
  | tr_val_abstract_ptr : ∀l p p',  
    tr_accesses gt p p' →  
    tr_val gt (val_abstract_ptr l p) (val_abstract_ptr l p')
```


Field grouping - values

And if we look at the struct grouping case:

```
Inductive tr_val (gt:group_tr) : val → val → Prop :=
| tr_val_struct_group : ∀Tt Tg s s' fg fs sg,
  gt = make_group_tr Tt fs Tg fg →
  fs ⊆ dom s →
  fg ∉ dom s →
  dom s' = (dom s \ - fs) ∪ {fg} →
  dom sg = fs →
  (* Contents of the grouped fields. *)
  s'[fg] = val_struct (typ_var Tg) sg →
  (∀ f,
    f ∈ dom sg →
    tr_val gt s[f] sg[f]) →
  (* Contents of the rest of the fields. *)
  (∀ f,
    f ∉ fs →
    f ∈ dom s →
    tr_val gt s[f] s'[f]) →
  tr_val gt (val_struct (typ_var Tt) s) (val_struct (typ_var Tt) s')
```

Field grouping - terms

Finally, we also need to change some of the terms. In particular, we look at the struct access case:

```
Inductive tr_trm (gt:group_tr) : trm → trm → Prop :=
| tr_trm_struct_access_group : ∀fs Tt fg Tg f op0 t1 op2 op1 t1',
  gt = make_group_tr Tt fs Tg fg →
  f ∈ fs →
  (* The access s.f *)
  op0 = prim_struct_access (typ_var Tt) f →
  (* The access s'.fg.f *)
  op1 = prim_struct_access (typ_var Tt) fg →
  op2 = prim_struct_access (typ_var Tg) f →
  tr_trm gt t1 t1' →
  tr_trm gt (trm_app op0 (t1::nil)) (trm_app op2 ((trm_app op1 (t1'::nil))::nil))
```

Field grouping - main theorem

In the end the theorem that we prove is:

Theorem `red_tr`: $\forall \text{gt LLC } C \ C' \ t \ t' \ v \ S \ S' \ m1 \ m1' \ m2,$
`red C LLC S m1 t m2 v` \rightarrow
`group_tr_ok gt C` \rightarrow
`tr_typdefctx gt C C'` \rightarrow
`tr_trm gt t t'` \rightarrow
`tr_stack gt S S'` \rightarrow
`tr_state gt m1 m1'` \rightarrow
`wf_typdefctx C` \rightarrow
`wf_trm C t` \rightarrow
`wf_stack C S` \rightarrow
`wf_state C m1` \rightarrow
`~is_error v` \rightarrow
 $\exists v' \ m2', \quad \text{tr_val gt } v \ v'$
 $\wedge \quad \text{tr_state gt } m2 \ m2'$
 $\wedge \quad \text{red } C' \ \text{LLC } S' \ m1' \ t' \ m2' \ v'.$

Array tiling

We need to know:

- The name of the array being changed.
- The new name for the tiles.
- The size of the tiles.

Similarly, we also define:

- `tiling_tr_ok`,
- `tr_typdefctx`,
- `tr_accesses`,
- `tr_val`,
- `tr_stack`,
- `tr_state` and
- `tr_trm`.

In this case, we change all the instances of `t[i]` to `t[i/N][i%N]`.

Array tiling - some specifics

A couple of important definitions used throughout:

Definition `nb_tiles (K I J:int) : Prop :=`
`J = I / K + If (I mod K = 0) then 0 else 1.`

Definition `tiled_indices (I J K i j k:int) : Prop :=`
`i = j * K + k`
`^ index I i`
`^ index J j`
`^ index K k.`

AoS to SoA

For this transformation, we need to know:

- The name of the struct being changed.
- The name of the array being changed.

TODO: Add something more here

High-to-low level transformation

The grammar is extended with:

```
Inductive val : Type :=  
  | val_concrete_ptr : loc → offset → val  
  | val_words : list word → val.
```

```
Inductive prim : Type :=  
  | prim_ll_get : typ → prim  
  | prim_ll_set : typ → prim  
  | prim_ll_new : typ → prim  
  | prim_ll_access : typ → prim.
```

There are two sides of this transformation:

- The memory.
- The programs.

High-to-low level transformation - OK

We need to ensure consistency between the type definition context (\mathcal{C}) and the low-level context (LLC). In particular:

- The type variable sizes in LLC match with the types in \mathcal{C} .
- The field offsets match with the order of the fields and the sizes of each of their types.

High-to-low level transformation - memory




Figure: High-level memory.




Figure: Low-level memory.

High-to-low level transformation - program

Details on the transformation of terms.

High-to-low level transformation - proof

Main theorem.

Project extent

what has been done and what hasn't quite and statistics

Future work

for instance functions etc, combining them. Code realisations...

Conclusion

conclusion

Verification of Data Layout Transformations

Ramon Fernández Mir

with Arthur Charguéraud

Inria

17/09/2018